

Retrieval-Augmented Generation (RAG) is a technique in natural language processing that combines retrieval-based methods with generative models.

It retrieves relevant documents from an external corpus and uses them as context for a language model to generate accurate and grounded answers.

This document is used as a knowledge base in a RAG pipeline implemented in native Python using:

- FAISS for similarity search
- SentenceTransformers for text embeddings
- Groq API for language generation
- PyMuPDF for PDF reading

Use this document to test a question-answering system based on your own documents!